

US Open Men 2013 Analysis

April 16, 2017

Business Understanding

This is a public, multivariate dataset concerning the US Men's Tennis Open in 2013. Our data is sourced from UCI Machine Learning Repository (Jauhari, Morankar, Fokoue)

(<https://archive.ics.uci.edu/ml/datasets/Tennis+Major+Tournament+Match+Statistics>).

It has 42 variables in all. Each observation represents a singles tennis match.

In tennis, a match is composed of two to three sets. To win a match, a player must win two sets. A set is composed of at least six games. To win a set, one must win at least six games, with a two game advantage over the games won by the opposing player. The scoring of games in tennis is a little odd. The score begins 0-0 or "love." After the first point is scored, the score is "love"-15. The scoring continues sequentially, "love"(0 points), 15(2 points), 30(3 points), 40(4 points). The game must be won by two points. If the score is 40-40, that is called 40-all or "deuce," the one player must score an additional two points consecutively to win the game.

The response variable which we will be attempting to predict will be FNL.1 (the final number of games won by player 1). In predicting this variable, some other variables become obsolete such as Result and STX.Y. Four other variables that we will not include in our analysis are WNR.X and UFE.X as no data populates those fields. Below is the complete list of variables and their uses.

Use	Variable	Type	Description
None	Player 1		Name of player 1
None	Player 2		Name of player 2
None	Result	Ordinal (0/1)	Referenced on player 1. If player 1 wins- result=1 (FNL.1>FNL.2)
Covariate	FSP.1	Real Number	First serve percentage for player 1
Covariate	FSW.1	Real Number	First serve won by player 1
Covariate	SSP.1	Real Number	Second serve percentage for player 1
Covariate	SSW.1	Real Number	Second serve won by player 1
Covariate	ACE.1	Numeric-Integer	Aces won by player 1
Covariate	DBF.1	Numeric-Integer	Double faults committed by player 1
None	WNR.1	Numeric	Winners earned by player 1
None	UFE.1	Numeric	Unforced errors committed by player 1
Covariate	BPC.1	Numeric	Break points created by player 1
Covariate	BPW.1	Numeric	Break points won by player 1
Covariate	NPA.1	Numeric	Net points attempted by player 1
Covariate	NPW.1	Numeric	Net points won by player 1
Covariate	TPW.1	Numeric	Total points won by player 1
None	ST1.1	Numeric-Integer	Set 1 Result for player 1
None	ST2.1	Numeric-Integer	Set 2 Result for player 1
None	ST3.1	Numeric-Integer	Set 3 Result for player 1
None	ST4.1	Numeric-Integer	Set 4 Result for player 1
None	ST5.1	Numeric-Integer	Set 5 Result for player 1
Response	FNL.1	Numeric-Integer	Final number of games won by player 1
Covariate	FSP.1	Real Number	First serve percentage for player 2
Covariate	FSW.2	Real Number	First serve won by player 2
Covariate	SSP.2	Real Number	Second service percentage for player 2
Covariate	SSW.2	Real Number	Second serve won by player 2
Covariate	ACE.2	Numeric-Integer	Aces won by player 2
Covariate	DBF.2	Numeric-Integer	Double faults committed by player 2

Use	Variable	Type	Description
None	WNR.2	Numeric	Winners earned by player 2
None	UFE.2	Numeric	Unforced errors committed by player 2
Covariate	BPC.2	Numeric	Break points created by player 2
Covariate	BPW.2	Numeric	Break points won by player 2
Covariate	NPA.2	Numeric	Net points attempted by player 2
Covariate	NPW.2	Numeric	Net points won by player 2
Covariate	TPW.2	Numeric	Total points won by player 2
None	ST2.1	Numeric-Integer	Set 1 Result for player 2
None	ST2.2	Numeric-Integer	Set 2 Result for player 2
None	ST2.3	Numeric-Integer	Set 3 Result for player 2
None	ST2.4	Numeric-Integer	Set 4 Result for player 2
None	ST2.5	Numeric-Integer	Set 5 Result for player 2
None	FNL.2	Numeric-Integer	Final number of games won by player 2
Covariate	Round	Numeric-Integer	Round of the tournament at which the game is played

Data Understanding

- We will use `dim` function to get the dimensions of the data
- `str` function to get the variable information
- `head` function to read first few observations
- `summary` function to get a summary of the data set

```
usopen <- read.csv("USOpen-men-2013.csv")
dim(usopen)
```

```
## [1] 126 42
```

```
str(usopen)
```

```
## 'data.frame': 126 obs. of 42 variables:
## $ Player1: Factor w/ 81 levels "Adrian Mannarino",...: 65 74 37 7 26 47 10 60 24 54 ...
## $ Player2: Factor w/ 79 levels "Adrian Mannarino",...: 52 2 30 17 23 9 18 35 5 6 ...
## $ Round : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Result : int 1 1 0 0 1 0 0 0 0 1 ...
## $ FNL1 : int 3 3 2 1 3 1 2 1 2 3 ...
## $ FNL2 : int 0 0 3 3 1 3 3 2 3 0 ...
## $ FSP.1 : int 63 61 55 52 58 59 53 51 58 51 ...
## $ FSW.1 : int 45 44 61 41 54 68 59 39 68 35 ...
## $ SSP.1 : int 37 39 45 48 42 41 47 49 42 49 ...
## $ SSW.1 : int 16 19 32 19 30 37 44 24 31 27 ...
## $ ACE.1 : int 7 3 11 13 21 20 8 17 15 4 ...
## $ DBF.1 : int 7 2 13 8 3 11 8 6 8 5 ...
## $ WNR.1 : logi NA NA NA NA NA NA ...
## $ UFE.1 : logi NA NA NA NA NA NA ...
## $ BPC.1 : int 5 4 5 2 5 1 10 0 5 9 ...
## $ BPW.1 : int 16 13 13 9 16 1 18 1 22 11 ...
## $ NPA.1 : int 18 NA NA NA NA 30 NA 5 NA NA ...
## $ NPW.1 : int 25 NA NA NA NA 42 NA 10 NA NA ...
## $ TPW.1 : int 106 99 149 97 148 133 183 82 159 117 ...
## $ ST1.1 : int 6 6 6 5 6 7 6 3 6 6 ...
## $ ST2.1 : int 6 6 3 6 6 2 3 7 3 7 ...
## $ ST3.1 : int 6 6 6 3 6 6 7 3 6 7 ...
## $ ST4.1 : int NA NA 6 0 6 6 7 NA 6 NA ...
```

```
## $ ST5.1 : int NA NA 1 NA NA NA 4 NA 4 NA ...
## $ FSP.2 : int 59 56 55 55 61 71 59 55 49 49 ...
## $ FSW.2 : int 37 37 66 52 63 72 65 42 58 35 ...
## $ SSP.2 : int 41 44 45 45 39 29 41 45 51 51 ...
## $ SSW.2 : int 17 15 27 27 30 26 38 28 43 20 ...
## $ ACE.2 : int 6 18 10 16 8 6 2 12 10 4 ...
## $ DBF.2 : int 4 8 9 3 2 1 7 2 12 3 ...
## $ WNR.2 : logi NA NA NA NA NA NA ...
## $ UFE.2 : logi NA NA NA NA NA NA ...
## $ BPC.2 : int 1 0 5 6 0 3 11 2 4 7 ...
## $ BPW.2 : int 3 1 15 9 3 11 26 9 14 12 ...
## $ NPA.2 : int 30 NA NA NA NA 30 NA 10 NA NA ...
## $ NPW.2 : int 40 NA NA NA NA 48 NA 10 NA NA ...
## $ TPW.2 : int 83 71 149 121 123 151 187 104 160 103 ...
## $ ST1.2 : int 3 3 7 7 7 6 7 6 3 4 ...
## $ ST2.2 : int 4 3 6 4 2 6 6 6 6 6 ...
## $ ST3.2 : int 2 4 2 6 3 7 6 6 1 5 ...
## $ ST4.2 : int NA NA 2 6 3 7 5 NA 7 NA ...
## $ ST5.2 : int NA NA 6 NA NA NA 6 NA 6 NA ...
```

```
head(usopen)
```

```
##           Player1      Player2 Round Result FNL1 FNL2 FSP.1 FSW.1
## 1   Richard Gasquet Michael Russell    1      1      3      0      63      45
## 2   Stephane Robert Albano Olivetti    1      1      3      0      61      44
## 3 Jan-Lennard Struff Guillaume Rufin    1      0      2      3      55      61
## 4       Aljaz Bedene Dmitry Tursunov    1      0      1      3      52      41
## 5   Feliciano Lopez  Florent Serra    1      1      3      1      58      54
## 6   Kenny De Schepper Bradley Klahn    1      0      1      3      59      68
##   SSP.1 SSW.1 ACE.1 DBF.1 WNR.1 UFE.1 BPC.1 BPW.1 NPA.1 NPW.1 TPW.1 ST1.1
## 1     37     16      7      7    NA    NA      5     16     18     25    106      6
## 2     39     19      3      2    NA    NA      4     13     NA     NA     99      6
## 3     45     32     11     13    NA    NA      5     13     NA     NA    149      6
## 4     48     19     13      8    NA    NA      2      9     NA     NA     97      5
## 5     42     30     21      3    NA    NA      5     16     NA     NA    148      6
## 6     41     37     20     11    NA    NA      1      1     30     42    133      7
##   ST2.1 ST3.1 ST4.1 ST5.1 FSP.2 FSW.2 SSP.2 SSW.2 ACE.2 DBF.2 WNR.2 UFE.2
## 1      6      6    NA    NA     59     37     41     17      6      4    NA    NA
## 2      6      6    NA    NA     56     37     44     15     18      8    NA    NA
## 3      3      6      6      1     55     66     45     27     10      9    NA    NA
## 4      6      3      0    NA     55     52     45     27     16      3    NA    NA
## 5      6      6      6    NA     61     63     39     30      8      2    NA    NA
## 6      2      6      6    NA     71     72     29     26      6      1    NA    NA
##   BPC.2 BPW.2 NPA.2 NPW.2 TPW.2 ST1.2 ST2.2 ST3.2 ST4.2 ST5.2
## 1      1      3     30     40     83      3      4      2    NA    NA
## 2      0      1    NA    NA     71      3      3      4    NA    NA
## 3      5     15    NA    NA    149      7      6      2      2      6
## 4      6      9    NA    NA    121      7      4      6      6    NA
## 5      0      3    NA    NA    123      7      2      3      3    NA
## 6      3     11     30     48    151      6      6      7      7    NA
```

```
summary(usopen)
```

```
##           Player1      Player2      Round      Result
## Novak Djokovic      : 6   Rafael Nadal      : 7   Min.      :1   Min.      :0.0000
```

```

## Richard Gasquet : 6 David Ferrer : 5 1st Qu.:1 1st Qu.:0.0000
## Andy Murray : 5 Mikhail Youzhny: 4 Median :1 Median :0.0000
## Roger Federer : 4 Milos Raonic : 4 Mean :1 Mean :0.4683
## Stanislas Wawrinka: 4 Tomas Berdych : 4 3rd Qu.:1 3rd Qu.:1.0000
## Tommy Robredo : 4 Denis Istomin : 3 Max. :1 Max. :1.0000
## (Other) :97 (Other) :99
## FNL1 FNL2 FSP.1 FSW.1
## Min. :0.000 Min. :0.000 Min. :38.00 Min. :14.00
## 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:53.00 1st Qu.:38.00
## Median :2.000 Median :3.000 Median :59.00 Median :47.00
## Mean :1.841 Mean :1.881 Mean :58.65 Mean :47.44
## 3rd Qu.:3.000 3rd Qu.:3.000 3rd Qu.:63.00 3rd Qu.:57.00
## Max. :3.000 Max. :3.000 Max. :75.00 Max. :87.00
##
## SSP.1 SSW.1 ACE.1 DBF.1
## Min. :25.00 Min. : 8.00 Min. : 0.000 Min. : 0.000
## 1st Qu.:37.00 1st Qu.:17.00 1st Qu.: 5.000 1st Qu.: 3.000
## Median :41.00 Median :23.00 Median : 8.000 Median : 5.000
## Mean :41.35 Mean :23.38 Mean : 8.508 Mean : 4.952
## 3rd Qu.:47.00 3rd Qu.:29.00 3rd Qu.:11.000 3rd Qu.: 7.000
## Max. :62.00 Max. :45.00 Max. :29.000 Max. :14.000
##
## WNR.1 UFE.1 BPC.1 BPW.1
## Mode:logical Mode:logical Min. : 0.000 Min. : 0.00
## NA's:126 NA's:126 1st Qu.: 2.000 1st Qu.: 7.00
## Median : 5.000 Median :10.50
## Mean : 4.198 Mean :10.26
## 3rd Qu.: 6.000 3rd Qu.:14.00
## Max. :10.000 Max. :23.00
##
## NPA.1 NPW.1 TPW.1 ST1.1
## Min. : 4.00 Min. : 4.00 Min. : 38.0 Min. :0.000
## 1st Qu.:12.00 1st Qu.:19.75 1st Qu.: 91.0 1st Qu.:4.000
## Median :17.00 Median :25.50 Median :111.5 Median :6.000
## Mean :18.28 Mean :27.86 Mean :112.9 Mean :4.968
## 3rd Qu.:25.00 3rd Qu.:36.25 3rd Qu.:137.0 3rd Qu.:6.000
## Max. :39.00 Max. :63.00 Max. :195.0 Max. :7.000
## NA's :38 NA's :38
## ST2.1 ST3.1 ST4.1 ST5.1
## Min. :1.000 Min. :0.000 Min. :0.000 Min. :0.00
## 1st Qu.:3.000 1st Qu.:3.000 1st Qu.:3.000 1st Qu.:4.00
## Median :6.000 Median :6.000 Median :6.000 Median :4.00
## Mean :4.889 Mean :4.696 Mean :4.821 Mean :4.52
## 3rd Qu.:6.000 3rd Qu.:6.000 3rd Qu.:6.000 3rd Qu.:6.00
## Max. :7.000 Max. :7.000 Max. :7.000 Max. :7.00
## NA's :1 NA's :59 NA's :101
## FSP.2 FSW.2 SSP.2 SSW.2
## Min. :44.00 Min. :10.00 Min. :16.00 Min. : 2.00
## 1st Qu.:55.00 1st Qu.:34.25 1st Qu.:37.00 1st Qu.:16.00
## Median :58.00 Median :45.50 Median :42.00 Median :23.00
## Mean :58.92 Mean :46.94 Mean :41.08 Mean :23.13
## 3rd Qu.:63.00 3rd Qu.:59.50 3rd Qu.:45.00 3rd Qu.:28.00
## Max. :84.00 Max. :90.00 Max. :56.00 Max. :48.00
##

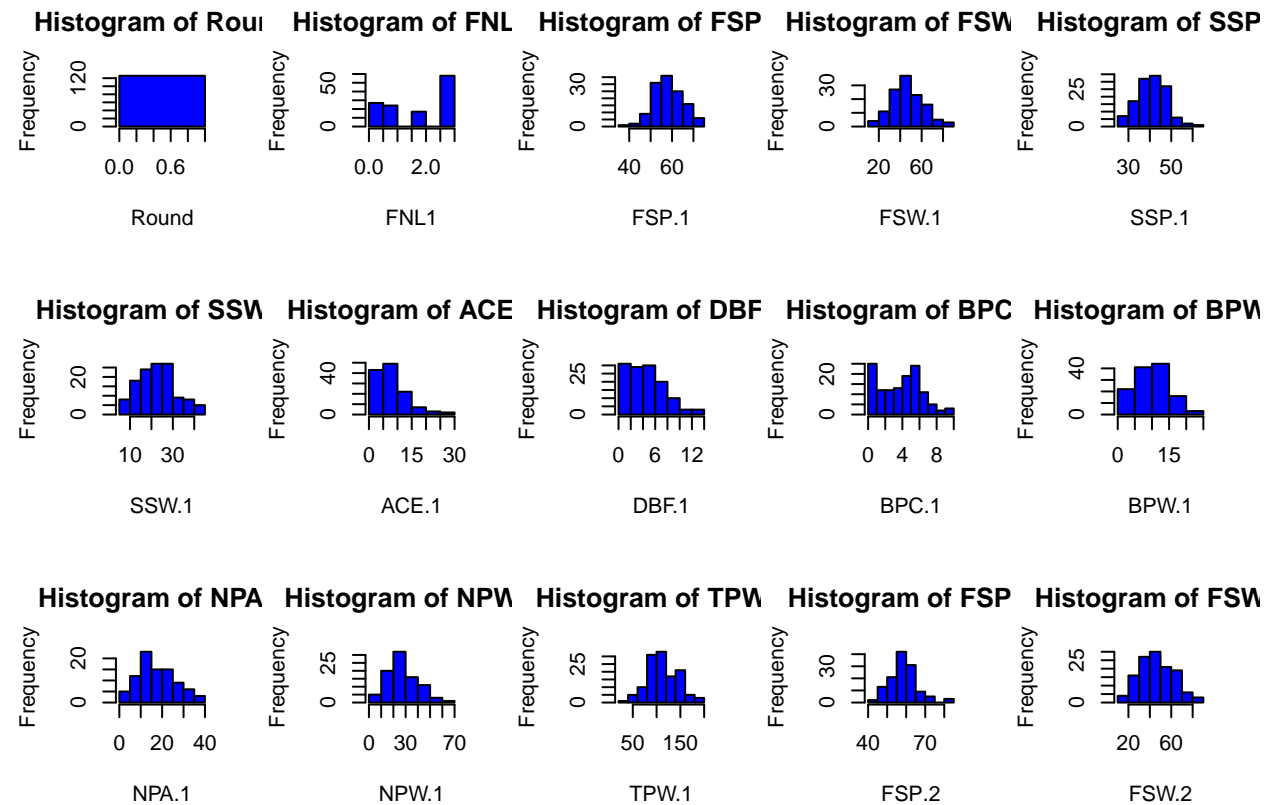
```

```
##      ACE.2      DBF.2      WNR.2      UFE.2
## Min.   : 0.000   Min.   : 0.000   Mode:logical   Mode:logical
## 1st Qu.: 4.250   1st Qu.: 3.000   NA's:126       NA's:126
## Median : 8.000   Median : 4.000
## Mean   : 9.262   Mean   : 4.595
## 3rd Qu.:11.000   3rd Qu.: 6.000
## Max.   :39.000   Max.   :15.000
##
##      BPC.2      BPW.2      NPA.2      NPW.2
## Min.   : 0.000   Min.   : 0.00   Min.   : 4.00   Min.   : 6.00
## 1st Qu.: 2.000   1st Qu.: 7.00   1st Qu.:12.00   1st Qu.:19.00
## Median : 4.000   Median :10.00   Median :18.50   Median :27.50
## Mean   : 4.087   Mean   :10.25   Mean   :19.84   Mean   :31.17
## 3rd Qu.: 6.000   3rd Qu.:13.75   3rd Qu.:26.00   3rd Qu.:41.00
## Max.   :11.000   Max.   :26.00   Max.   :48.00   Max.   :81.00
##
##                      NA's   :38   NA's   :38
##      TPW.2      ST1.2      ST2.2      ST3.2
## Min.   : 45.00   Min.   :0.000   Min.   :0.000   Min.   :0.000
## 1st Qu.: 87.25   1st Qu.:3.250   1st Qu.:3.000   1st Qu.:3.000
## Median :114.00   Median :6.000   Median :6.000   Median :6.000
## Mean   :113.18   Mean   :5.016   Mean   :4.516   Mean   :4.616
## 3rd Qu.:137.00   3rd Qu.:6.000   3rd Qu.:6.000   3rd Qu.:6.000
## Max.   :207.00   Max.   :7.000   Max.   :7.000   Max.   :7.000
##
##                      NA's   :1
##      ST4.2      ST5.2
## Min.   :0.000   Min.   :1.00
## 1st Qu.:4.000   1st Qu.:5.00
## Median :6.000   Median :6.00
## Mean   :5.015   Mean   :5.32
## 3rd Qu.:6.000   3rd Qu.:6.00
## Max.   :7.000   Max.   :7.00
## NA's   :59     NA's   :101
```

- Summary statistics
 - There are 126 observations with 42 variables
 - We read first few observations from the data set
 - WNR.1, WNR.2, UFE.2 and UFE.1 variables have no data
 - There are missing observations for ST4.1, ST5.1, NPA.2, NPW.2, ST3.2, ST4.2 and ST5.2 variables
- We use `hist` function to plot the histograms
- We use `plot` function to plot the density function
- As the variables are larger in number we avoid scatter plots at this moment


```
#Histograms
par(mfrow=c(3,5))
hist(Round, col="blue",xlab="Round")
hist(FNL1, col="blue",xlab="FNL1")
hist(FSP.1, col="blue",xlab="FSP.1")
hist(FSW.1, col="blue",xlab="FSW.1")
hist(SSP.1, col="blue",xlab="SSP.1")
hist(SSW.1, col="blue",xlab="SSW.1")
hist(ACE.1, col="blue",xlab="ACE.1")
hist(DBF.1, col="blue",xlab="DBF.1")
hist(BPC.1, col="blue",xlab="BPC.1")
hist(BPW.1, col="blue",xlab="BPW.1")
```

```
hist(NPA.1, col="blue",xlab="NPA.1")
hist(NPW.1, col="blue",xlab="NPW.1")
hist(TPW.1, col="blue",xlab="TPW.1")
hist(FSP.2, col="blue",xlab="FSP.2")
hist(FSW.2, col="blue",xlab="FSW.2")
```




```
par(mfrow=c(3,5))
hist(SSP.2, col="blue",xlab="SSP.2")
hist(SSW.2, col="blue",xlab="SSW.2")
hist(ACE.2, col="blue",xlab="ACE.2")
hist(DBF.2, col="blue",xlab="DBF.2")
hist(BPC.2, col="blue",xlab="BPC.2")
hist(BPW.2, col="blue",xlab="BPW.2")
hist(NPA.2, col="blue",xlab="NPA.2")
hist(NPW.2, col="blue",xlab="NPW.2")
hist(TPW.2, col="blue",xlab="TPW.2")

#Density Plot
par(mfrow=c(3,5))
```




```
plot(density(Round), col="blue",xlab="Round")
plot(density(FNL1), col="blue",xlab="FNL1")
plot(density(FSP.1), col="blue",xlab="FSP.1")
plot(density(FSW.1), col="blue",xlab="FSW.1")
plot(density(SSP.1), col="blue",xlab="SSP.1")
plot(density(SSW.1), col="blue",xlab="SSW.1")
plot(density(ACE.1), col="blue",xlab="ACE.1")
plot(density(DBF.1), col="blue",xlab="DBF.1")
plot(density(BPC.1), col="blue",xlab="BPC.1")
plot(density(BPW.1), col="blue",xlab="BPW.1")
plot(density(NPA.1,na.rm=T), col="blue",xlab="NPA.1")
plot(density(NPW.1,na.rm=T), col="blue",xlab="NPW.1")
plot(density(TPW.1), col="blue",xlab="TPW.1")
plot(density(FSP.2), col="blue",xlab="FSP.2")
plot(density(FSW.2), col="blue",xlab="FSW.2")
```


ensity.default(x = R
ensity.default(x = F
ensity.default(x = F
ensity.default(x = F
ensity.default(x = S



ensity.default(x = S
ensity.default(x = A
ensity.default(x = D
ensity.default(x = B
ensity.default(x = B




y.default(x = NPA.1,
y.default(x = NPW.1,
ensity.default(x = T
ensity.default(x = F
ensity.default(x = F




```
par(mfrow=c(3,5))
plot(density(SSP.2), col="blue",xlab="SSP.2")
plot(density(SSW.2), col="blue",xlab="SSW.2")
plot(density(ACE.2), col="blue",xlab="ACE.2")
plot(density(DBF.2), col="blue",xlab="DBF.2")
plot(density(BPC.2), col="blue",xlab="BPC.2")
plot(density(BPW.2), col="blue",xlab="BPW.2")
plot(density(NPA.2,na.rm=T), col="blue",xlab="NPA.2")
plot(density(NPW.2,na.rm=T), col="blue",xlab="NPW.2")
```


`ensity.default(x = S`
`ensity.default(x = S`
`ensity.default(x = A`
`ensity.default(x = D`
`ensity.default(x = B`



`ensity.default(x = B`
`y.default(x = NPA.2,`
`y.default(x = NPW.2,`



- Most of the data is numeric with little or no cleaning required. We can replace the missing values with zero (or mean value) to simplify the data modeling process.

Data Preparation

- We will not consider the variable Round as it is a constant and is not impacting the response or regressor variables as evident from the scatter plots
- We have chosen the response variable as FNL1 - Final number of games won by player 1
- Covariates for consideration - FSP.1, FSW.1, SSP.1, SSW.1, ACE.1, DBF.1, BPC.1, BPW.1, NPA.1, NPW.1, TPW.1, FSP.2, FSW.2, SSP.2, SSW.2, ACE.2, DBF.2, BPC.2, BPW.2, NPA.2, NPW.2, TPW.2
- We observe NA values in NPA and NPW variables and replace them with 0

```
usopen$NPA.1[is.na(usopen$NPA.1)] <- 0
usopen$NPW.1[is.na(usopen$NPW.1)] <- 0
usopen$NPW.2[is.na(usopen$NPW.2)] <- 0
usopen$NPA.2[is.na(usopen$NPA.2)] <- 0
```

- Rest of the data looks pretty clean

Modeling

- We will utilize multiple linear regression method for this model.
- Based on the correlation matrix and partial f tests we will decide on the final list of covariates and final number of observations

- Correlation matrix for the selected covariates

```
cor(usopen[, c( "FSP.1", "FSW.1", "SSP.1", "SSW.1", "ACE.1", "DBF.1", "BPC.1", "BPW.1",
"NPA.1", "NPW.1", "TPW.1", "FSP.2", "FSW.2", "SSP.2", "SSW.2", "ACE.2", "DBF.2", "BPC.2",
"BPW.2", "NPA.2", "NPW.2", "TPW.2"
)])
```

- We observe high correlation between NPA.1 and NPW.1, TPW.1 and FSW.1, NPA.2 and NPW.2, FSW.2 and TPW.2
- Other combinations of variables are also correlated
- There is a good probability that we may experience multicollinearity in our model
- We split the data set into testing (30%) and training data (70%)

```
# setting the seed to make the partition reproducible
set.seed(999)
index <-
  sample(seq_len(nrow(usopen)), size = floor(0.70 * nrow(usopen)))

usopen_train <- usopen[index,]
usopen_test <- usopen[-index, ]
```

- We now create model based on training data
- `summary` function is used to get the model details

```
model_usopen <-
  lm(
    FNL1 ~ FSP.1 + FSW.1 + SSP.1 + SSW.1 + ACE.1 + DBF.1 + BPC.1 + BPW.1 + NPA.1 +
      NPW.1 + TPW.1 + FSP.2 + FSW.2 + SSP.2 + SSW.2 + ACE.2 + DBF.2 + BPC.2 +
      BPW.2 + NPA.2 + NPW.2 + TPW.2,
    data = usopen_train
  )
summary(model_usopen)
```

```
##
## Call:
## lm(formula = FNL1 ~ FSP.1 + FSW.1 + SSP.1 + SSW.1 + ACE.1 + DBF.1 +
##      BPC.1 + BPW.1 + NPA.1 + NPW.1 + TPW.1 + FSP.2 + FSW.2 + SSP.2 +
##      SSW.2 + ACE.2 + DBF.2 + BPC.2 + BPW.2 + NPA.2 + NPW.2 + TPW.2,
##      data = usopen_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.97603 -0.30841  0.03371  0.33374  1.05906
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.262626   1.376240  -0.917  0.36220
## FSP.1         0.002551   0.015872   0.161  0.87280
## FSW.1        -0.002686   0.018716  -0.144  0.88630
## SSP.1           NA         NA       NA     NA
## SSW.1         0.007278   0.025829   0.282  0.77899
## ACE.1         0.005061   0.013210   0.383  0.70284
## DBF.1         0.024640   0.025376   0.971  0.33504
## BPC.1         0.116067   0.063603   1.825  0.07248 .
## BPW.1        -0.049449   0.026547  -1.863  0.06689 .
```

```
## NPA.1      -0.019090    0.030332   -0.629   0.53125
## NPW.1       0.007268    0.020210    0.360   0.72028
## TPW.1       0.061255    0.014343    4.271  6.29e-05 ***
## FSP.2       0.024471    0.015574    1.571   0.12083
## FSW.2      -0.009003    0.020432   -0.441   0.66090
## SSP.2       NA         NA         NA      NA
## SSW.2       0.027966    0.023308    1.200   0.23443
## ACE.2       0.021018    0.010839    1.939   0.05671 .
## DBF.2      -0.004163    0.029013   -0.143   0.88633
## BPC.2       0.086542    0.073086    1.184   0.24055
## BPW.2      -0.025209    0.026448   -0.953   0.34393
## NPA.2       0.041423    0.025749    1.609   0.11238
## NPW.2      -0.023605    0.017060   -1.384   0.17105
## TPW.2      -0.053608    0.016279   -3.293   0.00158 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4957 on 67 degrees of freedom
## Multiple R-squared:  0.8815, Adjusted R-squared:  0.8462
## F-statistic: 24.93 on 20 and 67 DF,  p-value: < 2.2e-16
```

- The null hypothesis is $H_0 : \beta_1 = \beta_2 = \dots = \beta_p = 0$
- That is, there is not a single predictor which can be considered statistically significant
- The alternate hypothesis is H_a : At least one β_j is not zero
- That is, there is at least predictor which can explain the change in resultant variable
- We reject null hypothesis when the p value is < 0.05
- From the model summary we observe that there are only 2 statistically significant variables (null hypothesis is rejected for these)
 - TPW.1
 - TPW.2
- The F-statistic is 24.93 and the corresponding p-value is significantly lower than 0.05 so we can conclude to reject that null hypothesis that no predictor variable explains the variability in the response variable
- The R^2 value is 0.8815
- The model explains 88.15% of the variability in FNL1
- We now do a partial f-test for the variables FSP.1, FSW.1, SSP.1, SSW.1, ACE.1, DBF.1, NPA.1, NPW.1, FSP.2, FSW.2, SSP.2, SSW.2, ACE.2, DBF.2, NPA.2, NPW.2 and BPC.2

```
model_usopen_p <-
  lm(FNL1 ~ BPC.1 + BPW.1 + TPW.1 + BPW.2 + TPW.2,
     data = usopen_train)
anova(model_usopen, model_usopen_p)
```

```
## Analysis of Variance Table
##
## Model 1: FNL1 ~ FSP.1 + FSW.1 + SSP.1 + SSW.1 + ACE.1 + DBF.1 + BPC.1 +
##      BPW.1 + NPA.1 + NPW.1 + TPW.1 + FSP.2 + FSW.2 + SSP.2 + SSW.2 +
##      ACE.2 + DBF.2 + BPC.2 + BPW.2 + NPA.2 + NPW.2 + TPW.2
## Model 2: FNL1 ~ BPC.1 + BPW.1 + TPW.1 + BPW.2 + TPW.2
##   Res.Df    RSS  Df Sum of Sq    F Pr(>F)
## 1      67 16.465
```

```
## 2      82 20.008 -15    -3.5425 0.961 0.5046
```

- We observe that p value is > 0.05 and therefore all these variables are not statistically significant
- We can now exclude these variables from our analysis
- We summarize our current model

```
summary(model_usopen_p)
```

```
##
## Call:
## lm(formula = FNL1 ~ BPC.1 + BPW.1 + TPW.1 + BPW.2 + TPW.2, data = usopen_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.94966 -0.29081  0.04009  0.31831  1.20486
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.409225   0.201126   2.035  0.04512 *
## BPC.1        0.163649   0.043440   3.767  0.00031 ***
## BPW.1       -0.055868   0.017387  -3.213  0.00188 **
## TPW.1        0.048078   0.005446   8.828 1.58e-13 ***
## BPW.2       -0.024529   0.015967  -1.536  0.12834
## TPW.2       -0.034274   0.005362  -6.392 9.40e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.494 on 82 degrees of freedom
## Multiple R-squared:  0.856, Adjusted R-squared:  0.8473
## F-statistic: 97.53 on 5 and 82 DF,  p-value: < 2.2e-16
```

- We run another partial f-test for the variable BPW.2

```
model_usopen_p_2 <-
  lm(FNL1 ~ BPC.1 + BPW.1 + TPW.1 + TPW.2,
     data = usopen_train)
anova(model_usopen_p, model_usopen_p_2)
```

```
## Analysis of Variance Table
##
## Model 1: FNL1 ~ BPC.1 + BPW.1 + TPW.1 + BPW.2 + TPW.2
## Model 2: FNL1 ~ BPC.1 + BPW.1 + TPW.1 + TPW.2
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      82 20.008
## 2      83 20.584 -1   -0.57582 2.3599 0.1283
```

- We again get a p-value > 0.05
- Hence BPW.2 is also not statistically significant
- We summarize our current model

```
summary(model_usopen_p_2)
```

```
##
## Call:
## lm(formula = FNL1 ~ BPC.1 + BPW.1 + TPW.1 + TPW.2, data = usopen_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -1.04081 -0.40456 0.08174 0.36555 1.07170
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.397609   0.202623   1.962 0.053078 .
## BPC.1        0.132087   0.038586   3.423 0.000964 ***
## BPW.1       -0.053149   0.017437  -3.048 0.003090 **
## TPW.1        0.053237   0.004322  12.317 < 2e-16 ***
## TPW.2       -0.040555   0.003496 -11.599 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.498 on 83 degrees of freedom
## Multiple R-squared:  0.8519, Adjusted R-squared:  0.8448
## F-statistic: 119.4 on 4 and 83 DF,  p-value: < 2.2e-16
```

- We now standardize the regression coefficients using unit normal scaling

```
usopen_train_standard = as.data.frame(apply(usopen_train[, c("FNL1", "BPC.1", "BPW.1", "TPW.1", "TPW.2"),
(x - mean(x)) / sd(x)
}))
```


- We now create the new model using standardized values

```
model_usopen_std <-
  lm(FNL1 ~ BPC.1 + BPW.1 + TPW.1 + TPW.2,
     data = usopen_train_standard)
summary(model_usopen_std)
```

```
##
## Call:
## lm(formula = FNL1 ~ BPC.1 + BPW.1 + TPW.1 + TPW.2, data = usopen_train_standard)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.82346 -0.32007  0.06467  0.28921  0.84790
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.475e-16  4.200e-02   0.000 1.000000
## BPC.1        2.462e-01  7.192e-02   3.423 0.000964 ***
## BPW.1       -2.148e-01  7.047e-02  -3.048 0.003090 **
## TPW.1        1.426e+00  1.158e-01  12.317 < 2e-16 ***
## TPW.2       -1.125e+00  9.700e-02 -11.599 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.394 on 83 degrees of freedom
## Multiple R-squared:  0.8519, Adjusted R-squared:  0.8448
## F-statistic: 119.4 on 4 and 83 DF,  p-value: < 2.2e-16
```

- To better visualize the coefficients we plot the barplot

```
barplot(model_usopen_std$coefficients)
```



- We observe that the most statistically significant variables are TPW.1 and TPW.2
- We check our model for multicollinearity
- We will use the vif function from car library to examine multicollinearity

```
library(car)
vif(model_usopen_std)
```

```
##      BPC.1      BPW.1      TPW.1      TPW.2
## 2.899012 2.782959 7.514762 5.273466
```

- We observe that there is multicollinearity in our model (as was expected)
- We find the correlation matrix of the variables

```
cor(usopen_train_standard[, c("FNL1", "BPC.1", "BPW.1", "TPW.1", "TPW.2")])
```

```
##           FNL1           BPC.1          BPW.1          TPW.1          TPW.2
## FNL1      1.00000000  0.75807021  0.5302779  0.4943759 -0.06581077
## BPC.1     0.75807021  1.00000000  0.6804967  0.4540390 -0.00928395
## BPW.1     0.53027788  0.68049669  1.0000000  0.6824489  0.35180657
## TPW.1     0.49437585  0.45403905  0.6824489  1.0000000  0.79731644
## TPW.2    -0.06581077 -0.00928395  0.3518066  0.7973164  1.00000000
```

- Since TPW.1 and TPW.2 are highly correlated, we keep only TPW.1 and recreate the model

```
model_usopen_std_2 <-
  lm(FNL1 ~ BPC.1 + BPW.1 + TPW.1 ,
```

```
data = usopen_train_standard)
summary(model_usopen_std_2)
```

```
##
## Call:
## lm(formula = FNL1 ~ BPC.1 + BPW.1 + TPW.1, data = usopen_train_standard)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.32305 -0.46140 -0.00508  0.45556  1.20729
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.619e-17  6.759e-02   0.000  1.00000
## BPC.1        7.449e-01  9.279e-02   8.028 5.34e-12 ***
## BPW.1       -1.557e-01  1.131e-01  -1.377  0.17229
## TPW.1        2.624e-01  9.302e-02   2.821  0.00597 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6341 on 84 degrees of freedom
## Multiple R-squared:  0.6118, Adjusted R-squared:  0.598
## F-statistic: 44.14 on 3 and 84 DF,  p-value: < 2.2e-16
```

- We now calculate the VIF to check for multicollinearity

```
library(car)
vif(model_usopen_std_2)
```

```
##      BPC.1      BPW.1      TPW.1
## 1.863158 2.768419 1.872437
```

- There is no multicollinearity on this model
- As there is one statistically insignificant variable in this model, we run partial F test again for the variable BPW.1

```
model_usopen_std_2p <-
  lm(FNL1 ~ BPC.1 + TPW.1 ,
      data = usopen_train_standard)
summary(model_usopen_std_2p)
```

```
##
## Call:
## lm(formula = FNL1 ~ BPC.1 + TPW.1, data = usopen_train_standard)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.41846 -0.44510  0.05935  0.48807  1.28546
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.064e-16  6.795e-02   0.000  1.0000
## BPC.1       6.722e-01  7.670e-02   8.764 1.62e-13 ***
## TPW.1       1.892e-01  7.670e-02   2.467  0.0156 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.6374 on 85 degrees of freedom
## Multiple R-squared:  0.6031, Adjusted R-squared:  0.5937
## F-statistic: 64.58 on 2 and 85 DF,  p-value: < 2.2e-16
```

```
library(car)
vif(model_usopen_std_2p)
```

```
##      BPC.1      TPW.1
## 1.259686 1.259686
```

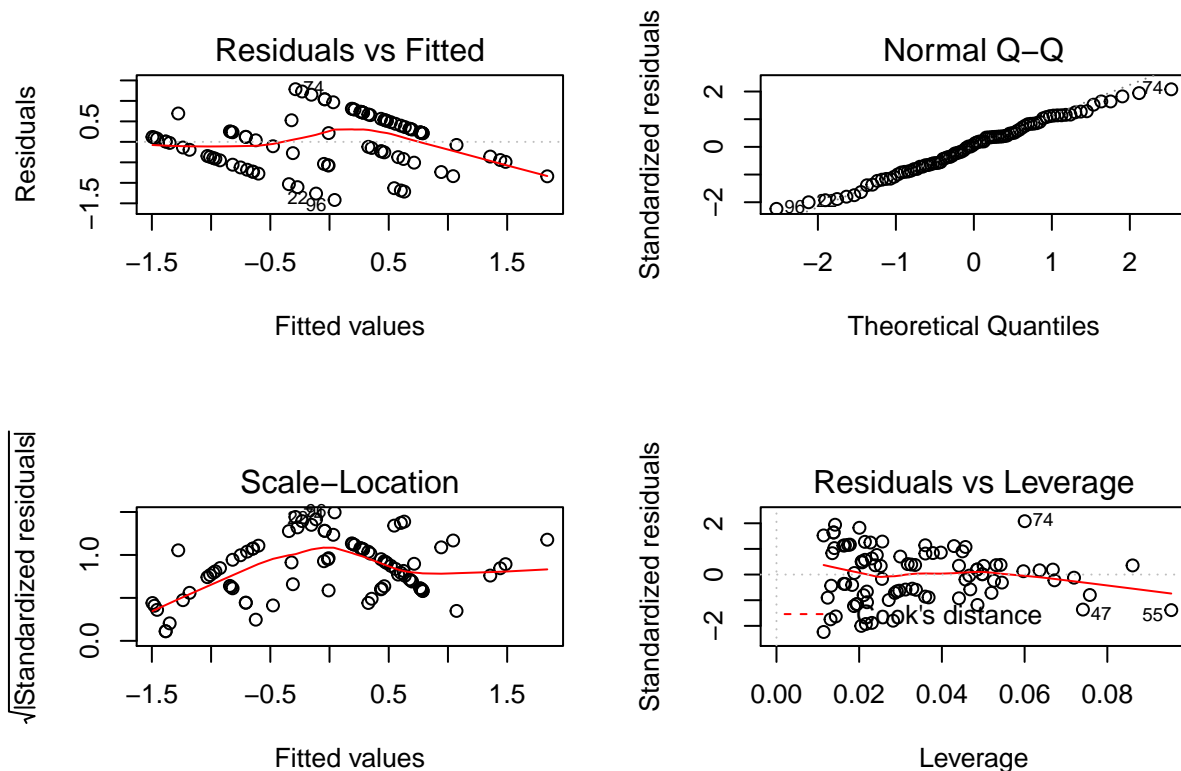
- This model does not have multicollinearity

```
anova(model_usopen_std_2, model_usopen_std_2p)
```

```
## Analysis of Variance Table
##
## Model 1: FNL1 ~ BPC.1 + BPW.1 + TPW.1
## Model 2: FNL1 ~ BPC.1 + TPW.1
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      84 33.770
## 2      85 34.532 -1   -0.76185 1.895 0.1723
```

- Using `anova` function we arrive to the conclusion that the variable BPW.1 is not statistically significant
- We now do the residual plots

```
par(mfrow = c(2, 2))
plot(model_usopen_std_2p)
```



Evaluation

- We calculate the mean square error for the models created

```
t <- predict(model_usopen_std, usopen_test)
mean((t - usopen_test[, c("FNL1")]) ** 2)
```

```
## [1] 1852.487
```

```
t <- predict(model_usopen_std_2, usopen_test)
mean((t - usopen_test[, c("FNL1")]) ** 2)
```

```
## [1] 969.4526
```

```
t <- predict(model_usopen_std_2p, usopen_test)
mean((t - usopen_test[, c("FNL1")]) ** 2)
```

```
## [1] 564.1643
```

- We observe that as we proceeded to refine our model the MSE kept on decreasing
- And we have the least MSE in our final model